

# SQL features in Trino

Taipei dbt meetup #35  
April 23, 2025

David Phillips

Co-creator of Trino and Presto



# What is Trino?

Fast, open source, SQL query engine that can query different data sources.



# Apache Iceberg



# Why Iceberg?

- Performance and scalability
- Hidden partitioning
- Safe schema evolution
- Time travel and rollback
- ACID transactions



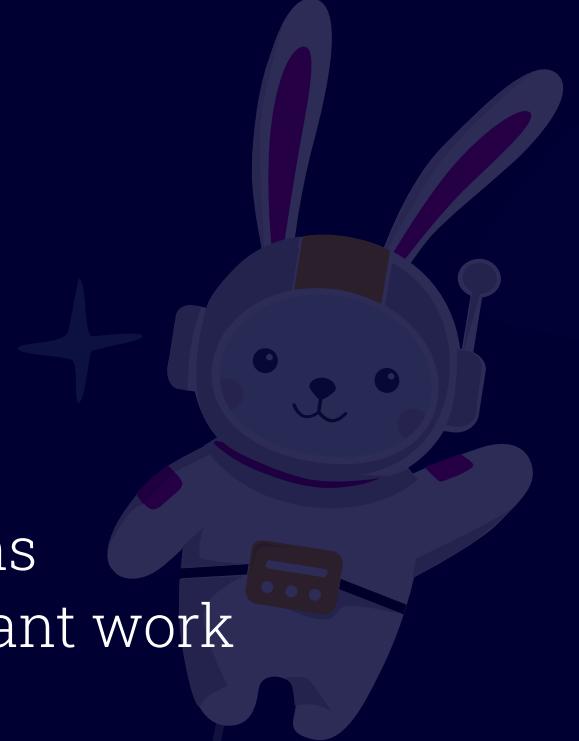
# Performance and scalability

- Fast scan planning
- Advanced filtering using metadata
- No file listing
- Supports huge tables



# Hive partitioning

- Partitions are explicit columns in the table
- Must use value-based partitioning
- Queries must explicitly filter on partition columns
- Partitioning **cannot be changed** without significant work



# Hidden partitioning

- Partitions are **implicit**: no partition columns in the table
- Example: partition by **day** or **hour** on a **timestamp** column
- Queries filter on the **original columns**
- Partition layouts can **evolve** over time



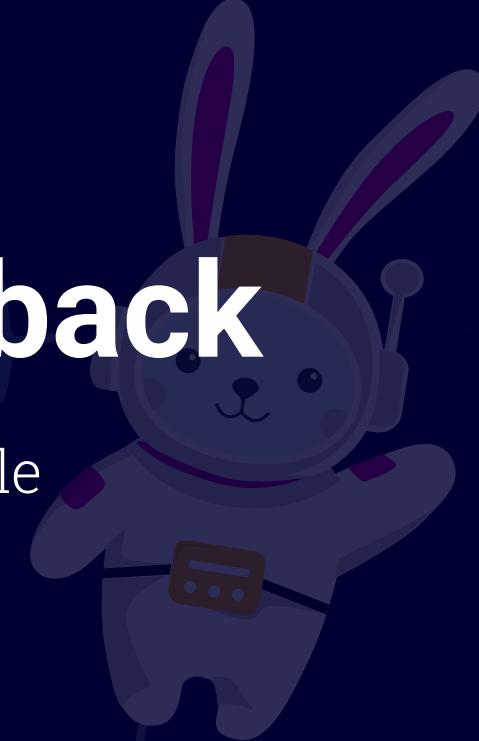
# Safe schema evolution

- Add, drop, or rename columns
- Reorder columns in a nested struct
- Widen column types
  - Convert from `int` to `bigint`



# Time travel and rollback

- View historical versions of a table
- Rollback to a previous version



# ACID transactions

- Data files are tracked in the table metadata
- Readers always see a **consistent view** of the data
- Concurrent writes are **safe** from conflicts

# What's new in Trino?



# Python user-defined functions

```
WITH FUNCTION double_it(x int)
    RETURNS int
    LANGUAGE PYTHON
    WITH (handler = 'twice')
    AS $$
def twice(a):
    return a * 2
$$
SELECT double_it(21); -- returns 42
```

```
CREATE FUNCTION double_it(x int)
    RETURNS int
    LANGUAGE PYTHON
    ...
```



# Spooling client protocol

- Much higher throughput for data transfer
- Faster query completion
- Reduces CPU usage on coordinator
- Direct download of data from spooling storage

# Other improvements

- Access control with [Open Policy Agent](#) or Apache Ranger
- Lineage tracking with [OpenLineage](#)
- Observability with [OpenTelemetry](#) and [OpenMetrics](#)

# SQL features in Trino



# Simple case expressions

```
SELECT n,  
       CASE n  
           WHEN 1 THEN 'one'  
           WHEN 2 THEN 'two'  
           ELSE 'many'  
       END AS name  
FROM (VALUES 1, 2, 3, 4) AS t (n);
```

n	name
1	one
2	two
3	many
4	many



# Searched case expressions

```
SELECT n,  
       CASE  
           WHEN n = 1 THEN 'aaa'  
           WHEN n IN (2, 3) THEN 'bbb'  
           ELSE 'ccc'  
       END AS name  
FROM (VALUES 1, 2, 3, 4) AS t (n);
```

n	name
1	aaa
2	bbb
3	bbb
4	ccc

# IF expression

```
SELECT format('I have %s cat%s.', n,
              IF(n = 1, '', 's')) AS text
FROM (VALUES 0, 1, 2, 3) AS t (n);
```

```
text
-----
I have 0 cats.
I have 1 cat.
I have 2 cats.
I have 3 cats.
```



# TRY expression

```
SELECT try(8 / 0) div_zero,  
       try(cast('abc' AS integer)) not_integer,  
       try(2000000000 + 2000000000) overflow;
```

TRY avoids these failures:

- Division by zero
- Cannot cast 'abc' to INT
- integer addition overflow: 2000000000 + 2000000000



# Conditional and conversion

```
SELECT try_cast('abc' AS integer); -- null  
  
SELECT coalesce(null, null, 'abc', '123'); -- 'abc'  
  
SELECT nullif('abc', 'abc'); -- null  
SELECT nullif('abc', '123'); -- 'abc'  
  
SELECT typeof('123'), typeof(123), typeof(123.0);
```



# Format function

```
SELECT format('pi = %.5f', pi()); -- 3.14159  
  
SELECT format('agent %03d', 7); -- agent 007  
  
SELECT format('$%,.2f', 1234567.89); -- $1,234,567.89  
  
SELECT format('%-7s,%7s', 'hello', 'world');  
-- hello , world  
  
SELECT format('%2$s %3$s %1$s', 'a', 'b', 'c'); -- b c a  
  
SELECT format('%1$tB %1$te, %1$tY', DATE '2001-07-04');  
-- July 4, 2001
```



# String comparison

LIKE operator with \_ and % placeholders

```
SELECT name FROM tpch.tiny.nation
WHERE name LIKE 'A%'
OR name LIKE '__N';
```



# Regular expressions

- `regexp_count`
- `regexp_extract`
- `regexp_like`
- `regexp_position`
- `regexp_replace`

# JSON functions

- json\_exists
- json\_query
- json\_value
- json\_array
- json\_object

Standard SQL JSON path language

# Date and timestamp functions

```
SELECT date_format(now(), '%d-%M-%Y %H'); -- MySQL
```

```
SELECT format_datetime(now(), 'y-M-dd'); -- Java
```

```
SELECT human_readable_seconds(37462);
```

```
SELECT format('%1$tB %1$te, %1$tY', DATE '2023-09-07');
```

```
SELECT extract(MONTH FROM TIMESTAMP '2023-11-20 18:10:00');
```

```
SELECT month(TIMESTAMP '2023-11-20 18:10:00');
```



# Many other functions

```
SELECT parse_data_size('2.3MB');
```

```
SELECT contains('10.0.0.0/8', IPADDRESS '11.255.255.255');
```

```
SELECT url_extract_host('https://trino.io/docs/');
```

```
SELECT uuid(); -- new random UUID with each call
```



# Structural data types

Many functions for arrays, maps, and rows

```
SELECT  
    ARRAY[1, 2, 3],  
    ARRAY['foo', 'bar', 'bazz'],  
    MAP(ARRAY['foo', 'bar'], ARRAY[1, 2]),  
    ROW(1, 2.0);''
```



# Associated max value

Find the clerk with the most expensive order:

```
SELECT max_by(clerk, totalprice) clerk,  
       max(totalprice) price  
  FROM orders;
```

# Pivoting using filtering

Order counts by priority

```
SELECT priority, count(*)  
FROM orders  
GROUP BY priority;
```

With separate columns per priority

```
SELECT  
    count(*) FILTER (WHERE priority = 'URGENT') AS urgent,  
    count(*) FILTER (WHERE priority = 'HIGH') AS high,  
    count(*) FILTER (WHERE priority = 'MEDIUM') AS medium,  
    count(*) FILTER (WHERE priority = 'LOW') AS low  
FROM orders;
```

trino

# ROLLUP with single column

```
SELECT  
    orderpriority,  
    count(*) AS orders  
FROM orders  
GROUP BY ROLLUP (orderpriority)  
ORDER BY orderpriority;
```



# ROLLUP with multiple columns

```
SELECT  
    linestatus, returnflag,  
    count(*) AS items  
FROM lineitem  
GROUP BY ROLLUP (linestatus, returnflag)  
ORDER BY linestatus, returnflag;
```



# CUBE

```
SELECT  
    linestatus, returnflag,  
    count(*) AS items  
FROM lineitem  
GROUP BY CUBE (linestatus, returnflag)  
ORDER BY linestatus, returnflag;
```



# GROUPING SETS

```
SELECT  
    linestatus, returnflag,  
    count(*) AS items  
FROM lineitem  
GROUP BY GROUPING SETS (  
    (linestatus),  
    (returnflag),  
    (linestatus, returnflag),  
    ()  
)  
ORDER BY linestatus, returnflag;
```



# Lambda expressions

Lambda expression with one input:

```
x -> x + 8
```

Lambda expression with two inputs:

```
(x, y) -> x + y
```



# Aggregating into an array

```
SELECT array_agg(name ORDER BY name DESC) names  
FROM region;
```

Note the ORDER BY clause



# Accessing array elements

```
SELECT a[2] AS second1,  
       element_at(a, 2) AS second2,  
       element_at(a, -2) AS second_from_last,  
       element_at(a, 99) AS bad  
  FROM (VALUES ARRAY[4, 5, 6, 7, 8]) AS t (a);
```

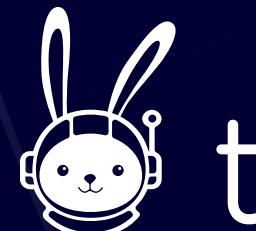


# Sorting arrays

```
SELECT array_sort(ARRAY['a', 'xyz', 'bb', 'abc', 'z', 'b']);
```

Optionally supply a comparator function

```
SELECT array_sort(ARRAY[3, 2, null, 5, null, 1, 2],  
                  -- sort null first with descending order  
                  (x, y) -> CASE WHEN x IS NULL THEN -1  
                           WHEN y IS NULL THEN 1  
                           WHEN x < y THEN 1  
                           WHEN x = y THEN 0  
                           ELSE -1 END);  
-- [null, null, 5, 3, 2, 2, 1]
```



# Matching elements

Do any, all, or none of the elements equal 8?

```
SELECT a,  
       any_match(a, e -> e = 8) AS any,  
       all_match(a, e -> e = 8) AS all,  
       none_match(a, e -> e = 8) AS none  
  FROM (VALUES ARRAY[4, 5, 6, 7, 8]) AS t (a);
```



# Filtering elements

Return data if function results in true

```
SELECT a,  
       filter(a, x -> x > 0) AS positive,  
       filter(a, x -> x IS NOT NULL) AS non_null  
FROM (VALUES ARRAY[5, -6, NULL, 7]) AS t (a);
```



trino

# Transforming elements

Convert each element with a function

```
SELECT a,  
       transform(a, x -> abs(x)) AS positive,  
       transform(a, x -> x * x) AS squared  
  FROM (VALUES ARRAY[5, -6, NULL, 7]) AS t (a);
```



# Converting arrays to strings

```
SELECT array_join(sequence(1, 5), '/'); -- 1/2/3/4/5
```



# Unnest an array

Split it up into multiple rows

```
SELECT name
FROM (
    VALUES ARRAY['cat', 'dog', 'mouse']
) AS t (a)
CROSS JOIN UNNEST(a) AS x (name);
```



# Unnest an array with ordinality

```
SELECT id, name
FROM (
    VALUES ARRAY['cat', 'dog', 'mouse']
) AS t (a)
CROSS JOIN UNNEST(a) WITH ORDINALITY AS x (name, id);
```



# Creating maps

From arrays of keys and values

```
SELECT map(ARRAY['x', 'y'], ARRAY[123, 456]);
```

From an array of entry rows

```
SELECT map_from_entries(ARRAY[('x', 123), ('y', 456)]);
```



# Accessing map elements

```
SELECT m,
       m['xyz'] AS xyz,
       element_at(m, 'abc') AS abc,
       element_at(m, 'bad') AS missing
  FROM (VALUES map_from_entries(
            ARRAY[('abc', 123), ('xyz', 456)])
      ) AS t (m);
```



# Unnesting a map

Split into multiple rows

```
SELECT key, value
FROM (VALUES map_from_entries(
          ARRAY[('abc', 123), ('xyz', 456)])
) AS t (m)
CROSS JOIN UNNEST(m) AS x (key, value);
```

# Thank you! Questions?

More Trino resources:

- Blog: [trino.io/blog](https://trino.io/blog)
- Documentation: [trino.io/docs](https://trino.io/docs)
- Community: [trino.io/community](https://trino.io/community)
- Join us on Slack: [trino.io/slack](https://trino.io/slack)

